

15-150 Fall 2013

Lecture 19

Stephen Brookes

today

parallel programming

- parallelism and functional style
- cost semantics
- Brent's Theorem and speed-ups
- sequences: an abstract type with efficient parallel operations

parallelism

- exploiting *multiple processors*
- evaluating *independent code simultaneously*
- low level implementation
 - *scheduling work onto processors*
- high level planning
 - designing code *abstractly*
 - without *baking in a schedule*

our approach

- design code *abstractly*
 - specify *behavioral correctness*
 - specify *asymptotic runtime (work, span)*
- reason about code *abstractly*
 - *independently* of schedule
 - *cost semantics and evaluation*
- You design the code
- The compiler schedules the work

functional benefits

- No side effects, so *evaluation order* doesn't affect *behavioral correctness*
- Can build *abstract types* that support efficient *parallel-friendly* operations
- Can use *work* and *span* to predict how *parallelizable* our code is
- Work and span are *independent* of scheduling details

caveat

- In practice, it's hard to achieve speed-up
- Current languages don't make it easy to implement good scheduling strategies
- Problems include:
 - scheduling overhead
 - locality of data (cache problems)
 - runtime sensitive to scheduling choices

why bother?

- It's good to learn to think *abstractly* first and figure out details later
 - Focus on *data dependencies* when you design your code
 - Our thesis: this approach to parallelism will *prevail*...
- (plus, I5-210 builds on these ideas...)

cost semantics

- We've already introduced *work* and *span*
- *Work* estimates the *sequential* running time on a *single* processor
- *Span* takes account of data dependency, estimates the *parallel* running time with *unlimited* processors
 - *critical path length* of computation

cost semantics

- We showed how to calculate *work* and *span* for *recursive functions* with ***recurrence relations***
- Now we introduce ***cost graphs***, another way to deal with work and span
- Cost graphs also allow us to talk about *schedules*...
- ... and the potential for *speed-up*


cost graphs

A cost graph is a series-parallel graph


- *directed graph, with source and sink*
- *nodes represent units of work* (constant time)
- *edges represent data dependencies*
- *branching indicates potential parallelism*

cost graphs

- a single node



sequential
composition



parallel
composition


work and span

of a cost graph


- The **work** is the *number of nodes*
- The **span** is the *length of the longest path from source to sink*

$$\mathbf{span}(G) \leq \mathbf{work}(G)$$

work




$$= \text{work } G_1 + \text{work } G_2 + c$$



$$= \text{work } G_1 + \text{work } G_2 + c$$


span

span




$$= \text{span } G_1 + \text{span } G_2 + c$$

span



$$= \max(\text{span } G_1, \text{span } G_2) + c$$

example



work = 11 (number of nodes)
span = 4 (longest path length)

using graphs


- Every *expression* can be given a *cost graph*
- Calculate the *work* and *span* using the graph
 - These are *asymptotically* the same as the work and span derived from *recurrence relations*

But what do work and span
tell us about the
actual running time?

scheduling

assign units of work to processors
respecting data dependency

- Work: number of steps taken by sequential scheduler on a single processor
- Span: number of steps taken by an optimal parallel scheduler with unlimited processes



(i) ① ② ⑥ ③ ④

(ii) ⑦ ⑤

(iii) ⑨ ⑧

(iv) ⑩

(v) ⑪


*an optimal
parallel schedule*

(5 rounds,
or 4 steps)

(uses 5 processors)

example

What if there are only 2 processors?



- (i) ① ②
- (ii) ③ ④
- (iii) ⑦ ⑤
- (iv) ⑥ ⑧
- (v) ⑨ ⑩
- (vi) ⑪

*a best schedule
for 2 processors*

(6 rounds,
5 steps)

2 processors

cannot do the job as fast as 5 processors (!)

Brent's Theorem


An expression with *work* **w** and *span* **s** can be evaluated on a **p**-processor machine in time $O(\max(\mathbf{w/p}, \mathbf{s}))$.

Optimal schedule using **p** processors:
Do (up to) **p** units of work each round
Total work to do is **w**
Needs at least **s** steps

What's the significance of the smallest **p** such that $\mathbf{w/p} \leq \mathbf{s}$?

Using more than this many processors won't yield any speed-up

example



work = 11
span = 4

$\min \{p \mid 11/p \leq 4\}$ is 3

- (i) 1 3 4
- (ii) 2 6 5
- (iii) 7 8
- (iv) 9 10
- (v) 11

*a best schedule
for 3 processors*

(5 rounds,
4 steps)

3 processors
can do the job as fast as 5(!)

next

- Exploiting parallelism in ML
- A signature for *parallel collections*
- *Cost analysis* of implementations
- *Cost benefits* of parallel algorithm design

sequences

signature SEQ =

sig

type 'a seq

exception Range

val tabulate : (int -> 'a) -> int -> 'a seq

val length : 'a seq -> int

val nth : int -> 'a seq -> 'a

val map : ('a -> 'b) -> 'a seq -> 'b seq

val reduce : (('a * 'a) -> 'a) -> 'a -> 'a seq -> 'a

val mapreduce : ('a -> 'b) -> 'b -> ('b * 'b -> 'b) -> 'a seq -> 'b

end

implementations

- Many ways to implement the signature
 - lists, balanced trees, arrays, ...
- For each one, can give a *cost analysis*
- There may be implementation *trade-offs*
 - *arrays: item access is $O(1)$*
 - *trees: item access is $O(\log n)$*

Seq : SEQ

- An abstract parameterized type of *sequences*
- Think of a sequence as a *parallel collection*
- With *parallel-friendly* operations
 - *constant-time* access to items
 - *efficient* map and reduce

We'll work today with an implementation

Seq : SEQ
based on *vectors*

notation

- We use math notation like

$$\langle v_1, \dots, v_n \rangle$$
$$\langle v_0, \dots, v_{n-1} \rangle$$
$$\langle \rangle$$

for sequence values

$$\langle 1, 2, 4, 8 \rangle : \text{int seq}$$

equality

- Two sequence values are (extensionally) *equal* iff they have the same length and their items are equal

$$\langle v_1, \dots, v_n \rangle = \langle u_1, \dots, u_m \rangle$$

if and only if

$$n = m \text{ and for all } i, v_i = u_i$$


operations

- For each operation in the signature SEQ we specify the (extensional) behavior of the operation implemented in Seq and discuss its cost semantics
- Other structures with the same signature may implement the operations with *different* work and span profile
- Learn to choose wisely!

tabulate

$\text{tabulate } f \ n = \langle f \ 0, \dots, f(n-1) \rangle$

- If G_i is cost graph for $f(i)$,
the cost graph for $\text{tabulate } f \ n$ is



work?
span?

If f is $O(1)$, the work for $\text{tabulate } f \ n$ is $O(n)$

If f is $O(1)$, the span for $\text{tabulate } f \ n$ is $O(1)$

examples

- `tabulate (fn x:int => x) 6`
- `tabulate (fn x:int => x*x) 6`
- `tabulate (fn _ => raise Range) 0`

length

$$\text{length } \langle v_1, \dots, v_n \rangle = n$$

- Work is $O(1)$
- Span is $O(1)$
- Cost graph is



Contrast: $\text{List.length } [v_1, \dots, v_n] = n$
work, span $O(n)$

nth

$\text{nth } i \langle v_0, \dots, v_{n-1} \rangle = v_i$ if $0 \leq i < n$
 $= \text{raise Range}$ otherwise

- Work is $O(1)$
- Span is $O(1)$
- Cost graph is




*Seq provides
constant-time access to items*

map

$$\text{map } f \langle v_1, \dots, v_n \rangle = \langle f v_1, \dots, f v_n \rangle$$

cost graph



- If f is constant time, $work \ O(n)$
 $span \ O(1)$

(contrast with List.map)

reduce

reduce should be used to *combine* a sequence using an *associative* function **g** with *identity element* **z**

- $g : t * t \rightarrow t$ is **associative** iff for all $x_1, x_2, x_3 : t$
 $g(x_1, g(x_2, x_3)) = g(g(x_1, x_2), x_3)$
- **z** is an *identity* for **g** iff for all $x : t$, $g(x, z) = x$
- When **g** is associative and **z** an identity we write

$$v_1 \text{ g } v_2 \text{ g } \dots \text{ g } v_n \text{ g } z$$

for the result of combining v_1, v_2, \dots, v_n, z using **g**

$$\begin{aligned} \text{reduce } g \text{ } z \langle v_1, \dots, v_n \rangle &= v_1 \text{ g } v_2 \text{ g } \dots \text{ g } v_n \text{ g } z \\ &= v_1 \text{ g } v_2 \text{ g } \dots \text{ g } v_n \end{aligned}$$

reduce

- When g is associative and z is an identity

$$\text{reduce } g \ z \ \langle v_1, \dots, v_n \rangle = v_1 \ g \ v_2 \ g \ \dots \ g \ v_n \ g \ z$$

- If g is constant time,


$$\text{reduce } g \ z \ \langle v_1, \dots, v_n \rangle$$

has *work* $O(n)$


and *span* $O(\log n)$

(Contrast with `foldr`, `foldl` on lists)

reduce (op +) 0 $\langle 1, 2, 3, 4 \rangle$




reduce (op +) 0 $\langle 1, 2, 3, 4, 5, 6, 7, 8 \rangle$



reduce cost

$$\text{reduce } g \text{ z } \langle v_1, \dots, v_{2n} \rangle = \\ g(\text{reduce } g \text{ z } \langle v_1, \dots, v_n \rangle, \text{reduce } g \text{ z } \langle v_{n+1}, \dots, v_{2n} \rangle)$$

$$G_{\langle 1, \dots, 2n \rangle} =$$



$$W(2n) = 2 * W(n) + c$$

$$S(2n) = S(n) + c$$

$$W(n) \text{ is } O(n)$$

$$S(n) \text{ is } O(\log_2 n)$$

mapreduce

- When g is associative and z is an identity,

$$\text{mapreduce } f \ z \ g \ \langle v_1, \dots, v_n \rangle = (f \ v_1) \ g \ \dots \ g \ (f \ v_n) \ g \ z$$

- When f, g are constant time,

$$\text{mapreduce } f \ z \ g \ \langle v_1, \dots, v_n \rangle$$

has *work* $O(n)$

and *span* $O(\log n)$

examples

```
fun sum (s : int seq) : int =  
    reduce (op +) 0 s
```

```
fun count (s : int seq seq) : int =  
    sum (map sum s)
```

analysis

```
fun sum (s : int seq) : int = reduce (op +) 0 s
```

```
fun count (s : int seq seq) : int = sum (map sum s)
```


- Let **s** be a value of type **int seq seq** consisting of *n rows*, each of length *n*
- What are the work and span for

count s ?

analysis


Let $s = \langle s_1, \dots, s_n \rangle$, $s_i = \langle x_{i1}, \dots, x_{in} \rangle$, $t_i = \text{sum } s_i$

For each i , $\text{sum } s_i = \text{reduce}(\text{op } +) 0 \langle x_{i1}, \dots, x_{in} \rangle$



work is $O(n)$
span is $O(\log n)$

$\text{map sum } s = \langle \text{sum } s_1, \dots, \text{sum } s_n \rangle$




work is $O(n^2)$
span is $O(\log n)$

analysis

Let $t_i = \text{sum } s_i$

$\text{count } s = \text{sum } \langle t_1, \dots, t_n \rangle$



work is $O(n^2)$
span is $O(\log n)$